

Exercises for the lecture
Fundamentals of Simulation Methods
WS 2015/16

Lecturer: Volker Springel

Exercise sheet 1 (*due date: Oct 23, 2015, 11am*)

Issues of floating point arithmetic

1) Machine epsilon

Write a computer program in C, C++, or Python¹ that experimentally determines the machine epsilon ϵ_m , i.e. the smallest number ϵ_m such $1 + \epsilon_m$ still evaluates to something different from 1, for the following data types:

- (a) float
- (b) double
- (c) long double

2) Pitfalls of floating point arithmetic

Consider the following numbers:

```
double a = 1.0e17;
double b = -1.0e17;
double c = 1.0;
double x = (a + b) + c;
double y = a + (b + c);
```

Calculate the results for x and y. Which one is correct, if any? Explain, why the law of associativity is here broken.

3) Pitfalls of floating point representation

Consider the following C/C++ code:

¹Special note for Python: Floating point numbers in python have a fixed default precision. In order to enforce a precision like you do in C/C++, you have to use the numpy module. You can create a 32 bit variable like this:

```
a = np.float32(2.)
```

You have to encapsulate numerical constants in your operations as well:

```
b = a * np.float32(1.5)
```

You can check the type of a variable with "type(a)" (This should give "numpy.float32").

Bonus question for python users: What is the default machine precision of python floats?

```

float    x = 0.01;
double   y = x;
double   z = 0.01;

int      i = x * 10000;
int      j = y * 10000;
int      k = z * 10000;

printf("%d %d %d\n", i, j, k);

```

which prints out three integer numbers.

- (a) Explain why the numbers are not all equal.
- (b) Determine the rational number n/m , where n and m are natural numbers, that is represented by the single-precision IEEE-754 floating point variable x in the above example. Note: This number is not $1/100$.

4) Packing of numbers

Estimate how many numbers there are in the interval between 1.0 and 2.0, and in between the interval of 511.0 to 512.0, for IEEE-754

- (a) single precision
- (b) and double precision

numbers.

5) Summing a long list of numbers

On the lecture's moodle-site, you'll find a binary file *numbers.dat* (8 MB). This contains first a 32-bit integer number that gives the number of double-precision values stored in the file (1 million in the provided example), followed by the numbers themselves. (The file is in little-endian. If you happen to work on a big endian processor, which is unlikely these days, you need to swap the endianness.)

Write a read-statement for these numbers, and then try to sum them up, using different approaches.

- (a) First, sum the numbers with a simple loop, sequentially from the beginning to the end. Write down the result.
- (b) Next, sum them from the end to the beginning, reversing the initial direction of the loop. Write down the result.
- (c) Sort the numbers by their magnitude, and sum them from small to large. What do you get now?

- (d) Repeat the last experiment by using a summation variable of type `long double`. Do you think the obtained result is trustworthy and correct? (Optional: Try to get real ‘quad-double’ precision to work where the machine epsilon is of the order of 10^{-35} or so. Check the manual of your compiler what `long double` actually stands for... and about how you can switch on real 128-bit floating point precision emulated in software.)
- (e) Write a program that *exactly* sums the list of numbers using the GMP library (GNU big number library) that allows arbitrary precision floating point accuracy. What do you get?